

```
In [1]: #import the Standard Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #Load the data using pandas read functions
data=pd.read_csv('C://Users//Admin//Downloads//archive (3)//Sleep_health_and_lifestyle.csv')
data.head()
```

Out[2]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Health Risk
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	

```
In [3]: # Drop the person id columns
data=data.drop('Person ID',axis=1)
```

```
In [4]: #Checking the data shape
print(f"The data set contains {data.shape[1]} columns and rows are {data.shape[0]}")
```

The data set contains 12 columns and rows are 374

```
In [5]: # ALL the columns information
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Gender            374 non-null    object  
 1   Age               374 non-null    int64   
 2   Occupation        374 non-null    object  
 3   Sleep Duration    374 non-null    float64 
 4   Quality of Sleep  374 non-null    int64   
 5   Physical Activity Level 374 non-null    int64  
 6   Stress Level      374 non-null    int64   
 7   BMI Category      374 non-null    object  
 8   Blood Pressure    374 non-null    object  
 9   Heart Rate         374 non-null    int64   
 10  Daily Steps       374 non-null    int64  
 11  Sleep Disorder    374 non-null    object  
dtypes: float64(1), int64(6), object(5)
memory usage: 35.2+ KB
```

In [6]: `#Checking the Null values  
data.isna().sum()/len(data.shape)*100`

Out[6]:

Gender	0.0
Age	0.0
Occupation	0.0
Sleep Duration	0.0
Quality of Sleep	0.0
Physical Activity Level	0.0
Stress Level	0.0
BMI Category	0.0
Blood Pressure	0.0
Heart Rate	0.0
Daily Steps	0.0
Sleep Disorder	0.0
dtype: float64	


In [7]: `data.describe().style.background_gradient(cmap='nipy_spectral')`

Out[7]:

	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	Daily Steps
<b>count</b>	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
<b>mean</b>	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775	6816.844920
<b>std</b>	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676	1617.915679
<b>min</b>	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3000.000000
<b>25%</b>	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000	5600.000000
<b>50%</b>	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000	7000.000000
<b>75%</b>	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000	8000.000000
<b>max</b>	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000	10000.000000

In [8]: `#Visualize the Correlation map  
plt.figure(figsize=(8,6))  
sns.heatmap(data.corr(),cmap='vlag_r',annot=True,fmt='.2f', linewidths=1)`

```
plt.title("Visualize the correlation map")
plt.show()
```



# Explore Data Analysis

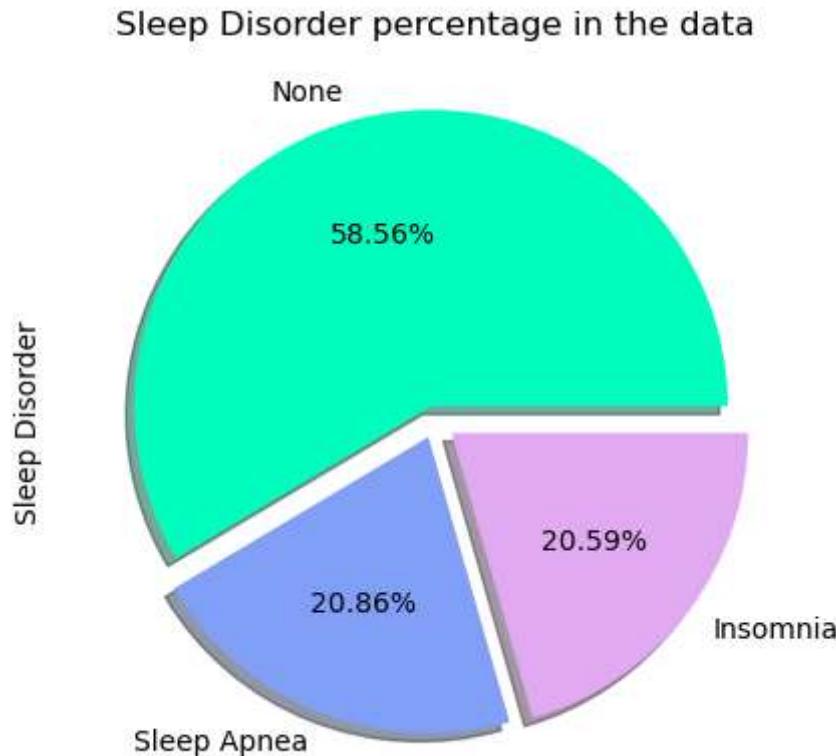
## Questions Asked for data:

- Sleep Disorder Percentage:
- Gender Percentage in the Data using a pie chart:
- Distribution of Age using a histogram
- Determine the highest occupation in the dataset.
- Analyze the distribution of sleep duration based on gender.
- Visualize the average sleep duration across different occupations using a bar chart.
- Explore the relationship between average sleep duration and BMI category.
- Identify the dominant occupation within the male category.
- Find the Average Heart with Bmi category

```
In [9]: # Find the Sleep disorder percentage
data['Sleep Disorder'].value_counts(normalize=True).sort_values(ascending=False)\n.plot(kind='pie',title="Sleep Disorder percentage in the data",explode=[0.05,0.06,0.07])
```


```
labels=['None', 'Sleep Apnea', 'Insomnia'],
colors=['#00FFBF', '#819FF7', '#E2A9F3'],
autopct='%.2f%%',
shadow=True)
```

Out[9]: <AxesSubplot:title={'center':'Sleep Disorder percentage in the data'}, ylabel='Sleep Disorder'>



In [10]: # Find the Gender Percentage using data  
data['Gender'].value\_counts(normalize=True).sort\_values(ascending=False)\\  
.plot(kind='pie', title="Gender Percentage in the data", explode=[0.05, 0.02],  
 labels=['Male', 'Female'],  
 colors=['#64FE2E', '#F5DA81'],  
 autopct='%.2f%%',  
 shadow=True)  
plt.show()

### Gender Percentage in the data

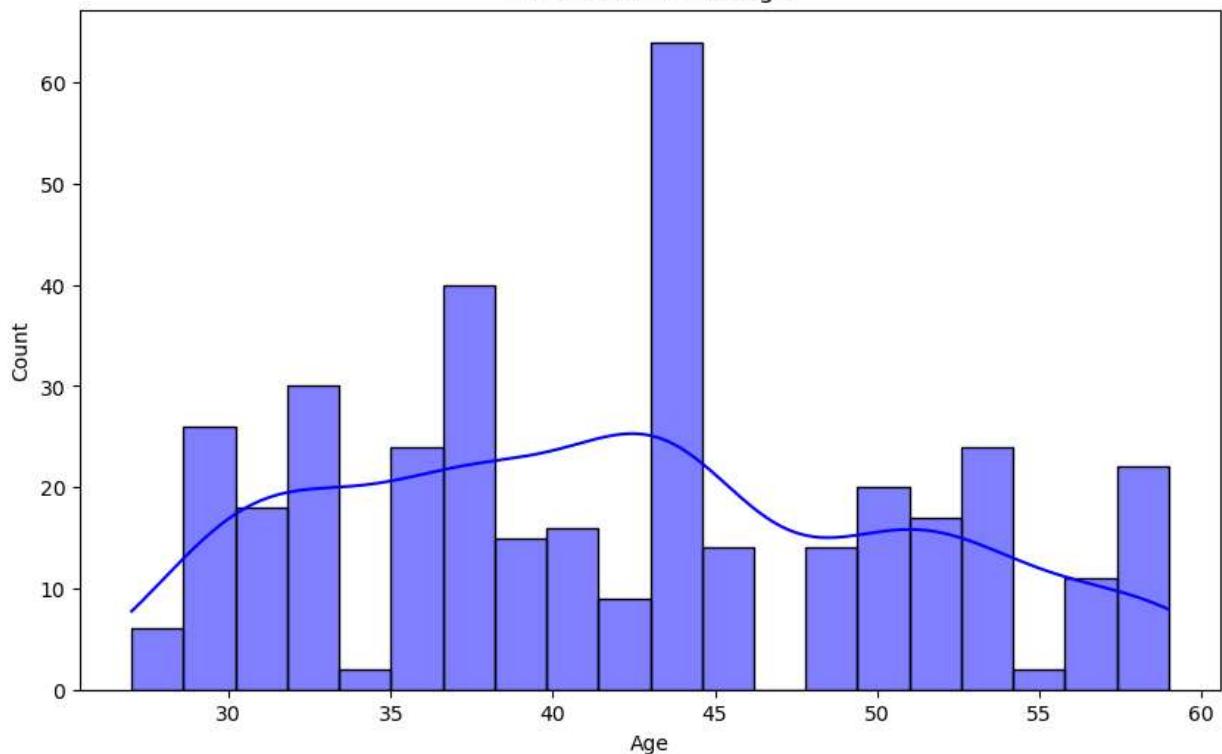


#### Observations:

- From the above two pie charts, we can observe several pieces of information.
- Firstly, the highest percentage in the sleep disorder pie chart is "None," indicating that a significant portion of the data does not have reported sleep disorders. The second-highest sleep disorder category is "Sleep Apnea."
- In the second pie chart depicting gender percentages, the male percentage is higher compared to the female percentage.


```
In [11]: #Distribution of the age columns
plt.figure(figsize=(10,6))
sns.histplot(data=data,x='Age',stat='count',bins=20,kde=True,color='b')
plt.title("Distribution of the Age")
plt.show()
```

## Distribution of the Age



```
In [12]: #Count of the occupation in the data using the bar plot
data['Occupation'].value_counts().sort_values(ascending=False) \
.plot(kind='bar',title="Visualize the Occupation",figsize=(8,6),color=['#0B0B61','#00F0F0']) \
plt.xlabel('Occupation') \
plt.ylabel('Count of Values') \
plt.show()
```

### Visualize the Occupation




### Observations:

- From the above two charts, one being a histogram and the other a bar chart, we can observe the following patterns:
- In the age histogram, there is a noticeable repetition of the age 43, indicating that a significant number of individuals in the dataset fall within this age range.
- In the bar chart, we can determine that the most demanding job based on the dataset is "Nurse," while the least demanding job is "Manager."


```
In [13]: #Visualize the Sleep Duration with Gender wise
plt.figure(figsize=(8,6))
sns.histplot(data=data,x='Sleep Duration',bins=30,kde=True,hue='Gender')
plt.title("Visualize the Sleep Duration with Gender wise")
plt.xlabel("Sleep Duration")
plt.ylabel("Count")
plt.show()
```

## Visualize the Sleep Duration with Gender wise



```
In [14]: #Find the Average Sleep time of the Different Occupations  
data.groupby('Occupation')['Quality of Sleep'].mean().sort_values(ascending=False)\  
.plot(kind='bar',title="Average Sleep time of different Occupation",figsize=(8,6),color  
plt.xlabel("Occupation")  
plt.ylabel("Average Sleep time ")  
plt.show()
```

Average Sleep time of different Occupation



## Observations:

- From the two charts, we can analyze the average sleep duration based on gender and the quality of sleep based on different occupations.
- According to the histogram chart, the average sleep duration for males is approximately 7.7 hours, while for females, it is around 8.5 hours.
- In the bar chart, it is evident that the occupation "Engineer" has the highest sleep quality among the different roles, while the occupation "Sales Representative" has the lowest sleep quality.

```
In [15]: #Some interesting questions asked in the data
#Gender imbalance
imbalanced_data=data['Gender'].value_counts()[0]/data['Gender'].value_counts()[1]
print('\nGender Imbalance',imbalanced_data)
#find the which is the dominate_occupation in the data
```

```

dominate_occupation=data['Occupation'].value_counts().idxmax()
print("\nDominate_occupation",dominate_occupation)
#find the Least demanding job
least_demanding_job=data['Occupation'].value_counts().idxmin()
print("\nThe Least demanding_job",least_demanding_job)
# find the top 5 stress level
top_5_stress_level=data['Stress Level'].value_counts().nlargest(5)
print("\nCommon stress pattern",top_5_stress_level)
#find the age range in the data
age_range=(data['Age'].min(),data['Age'].max())
print("\nAge Range",age_range)
# find the daily steps in the data
daily_steps=(data['Daily Steps'].min(),data['Daily Steps'].max())
print("\nRange of the Daily_steps",daily_steps)
# find the skewness of the sleep durations
sleep_quality_skewness=data['Sleep Duration'].value_counts().skew()
print("\nSleep Quality Skewness:", sleep_quality_skewness)

```

Gender Imbalance 1.0216216216216216

Dominate\_occupation Nurse

The Least demanding\_job Manager

```

Common stress pattern 3      71
8      70
4      70
5      67
7      50
Name: Stress Level, dtype: int64

```

Age Range (27, 59)

Range of the Daily\_steps (3000, 10000)

Sleep Quality Skewness: 0.7855254005718885

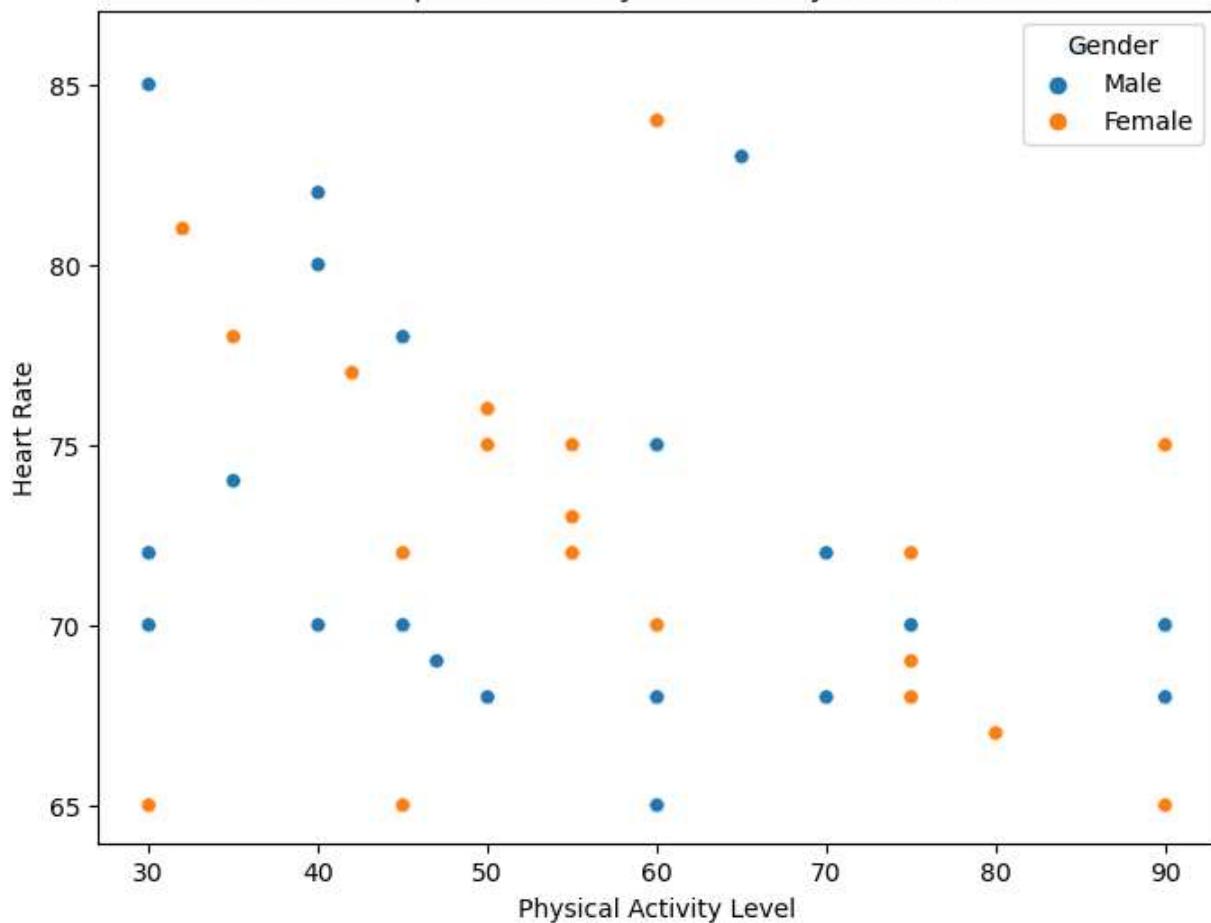
In [16]:

```

#Visualize the Physical Activity with Heart Rate on Gender
plt.figure(figsize=(8,6))
sns.scatterplot(data=data,x='Physical Activity Level',y='Heart Rate',hue='Gender')
plt.title("Relationship between Physical Activity and Heart Rate")
plt.show()

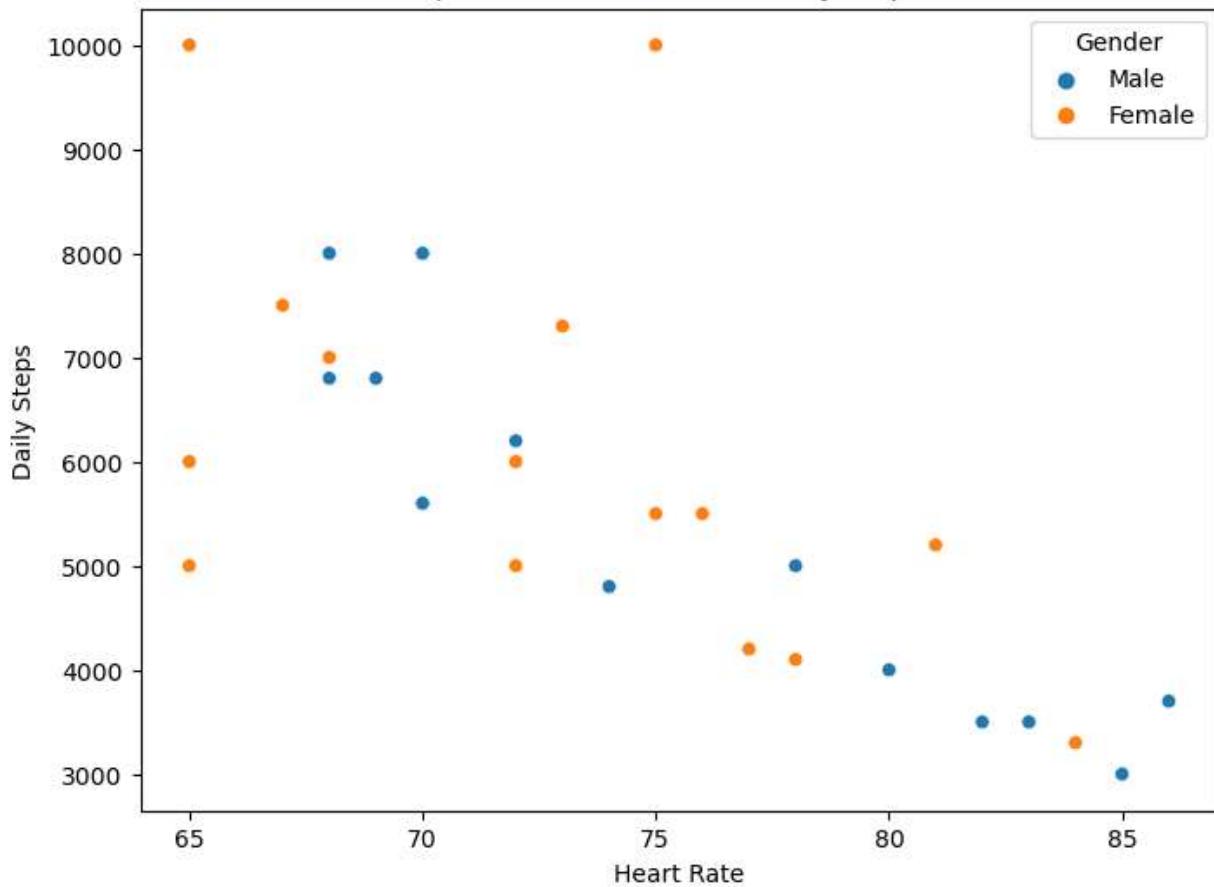
```

## Relationship between Physical Activity and Heart Rate



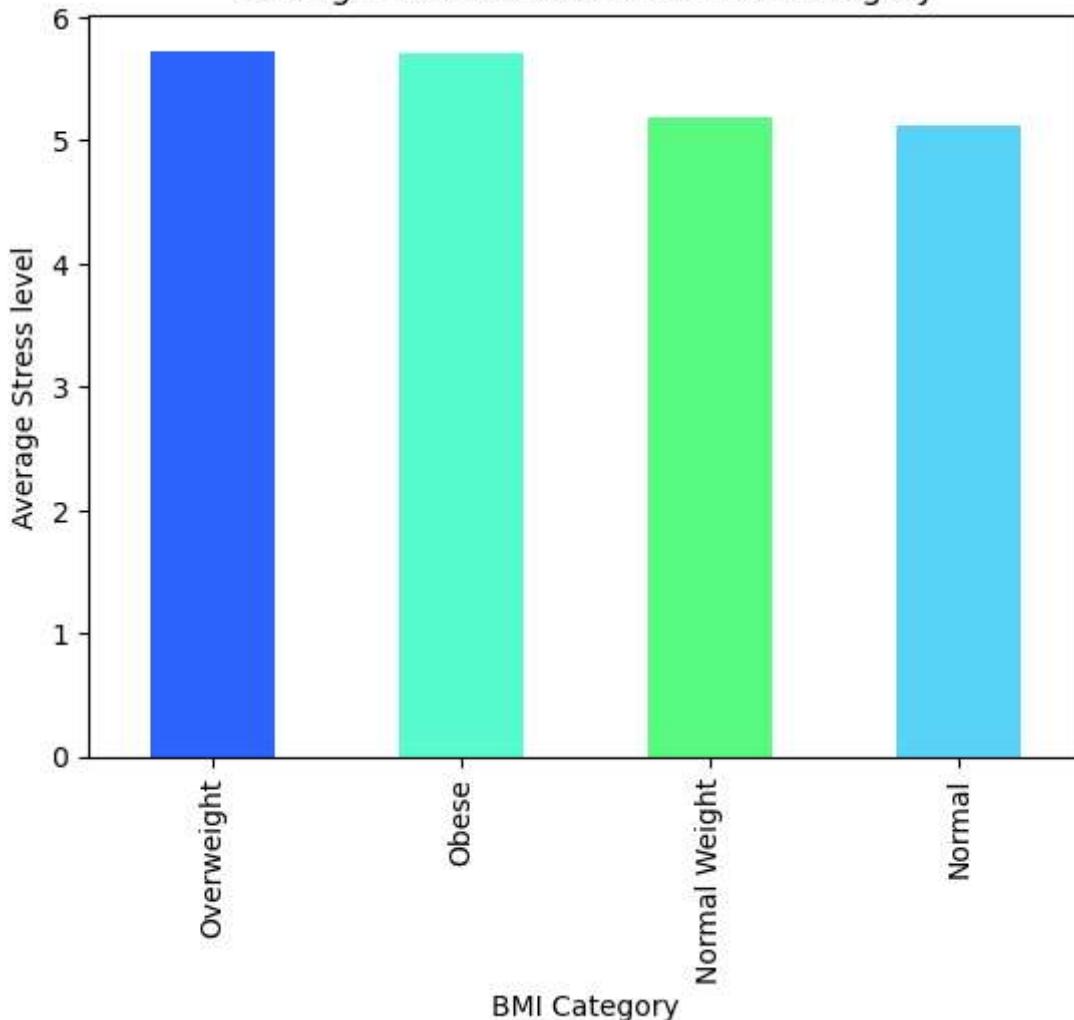
```
In [17]: # Find the Relationship between Heart Rate with Daily steps on gender
plt.figure(figsize=(8,6))
sns.scatterplot(data=data,x='Heart Rate',y='Daily Steps',hue='Gender')
plt.title("Relationship with Heart Rate and Daily steps on Gender")
plt.show()
```

## Relationship with Heart Rate and Daily steps on Gender



```
In [18]: #Find the average Stress Level the BMI Category wise  
data.groupby('BMI Category')['Stress Level'].mean().sort_values(ascending=False)\  
.plot(kind='bar',title='Average Strees level for the BMI Category',color=['#2E64FE','#  
plt.xlabel("BMI Category")  
plt.ylabel("Average Stress level")  
plt.show()
```

### Average Stress level for the BMI Category

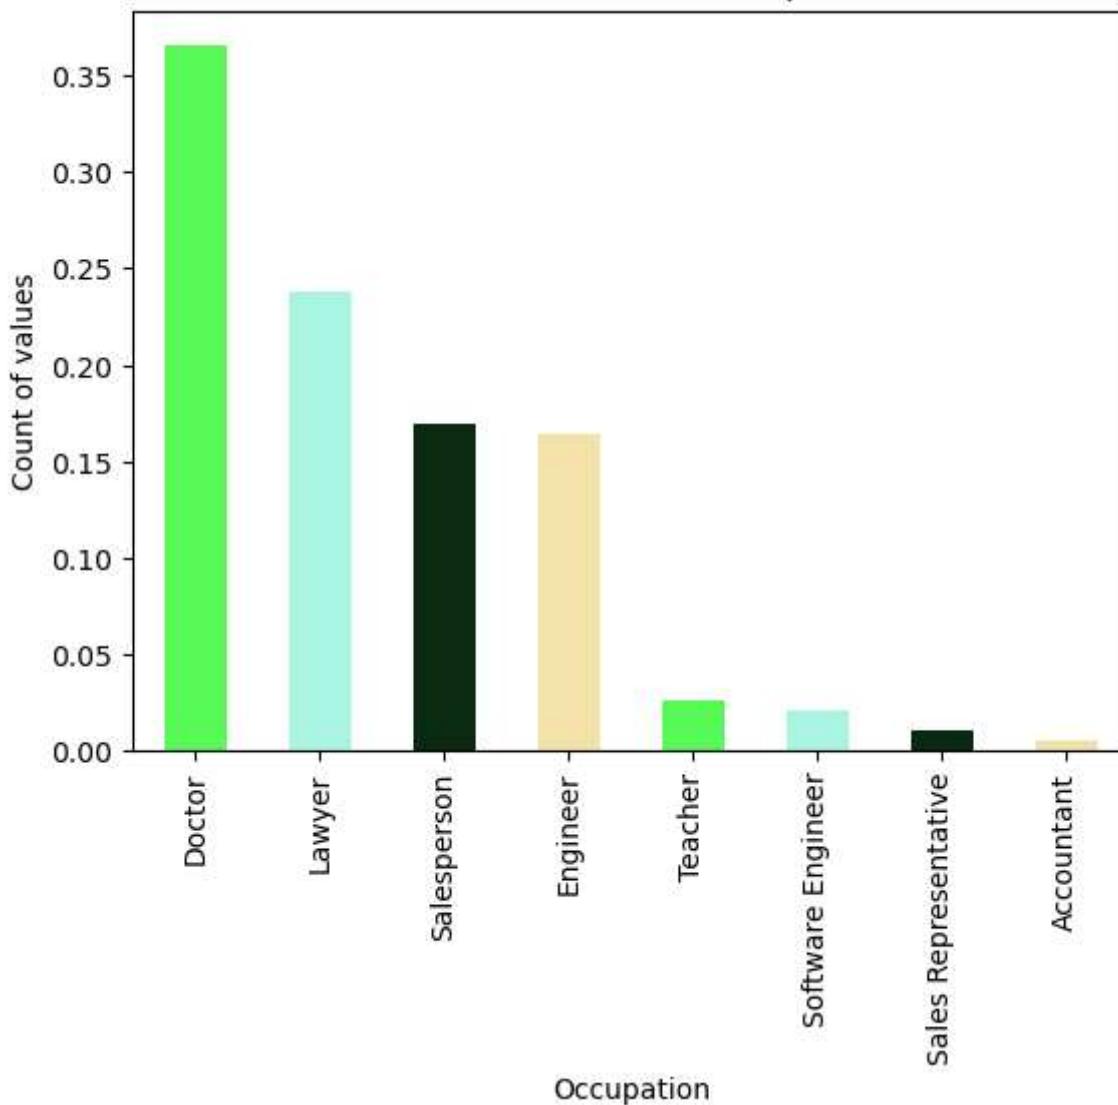


```
In [19]: #Find the averages daily steps with BMI Category on Gender using pivot_table
pd.pivot_table(data=data,index='Gender',values='Daily Steps',columns='BMI Category',ag
```

	Normal	Normal Weight	Obese	Overweight
Gender				
Female	6000.000000	6578.571429	3300.000000	7416.037736
Male	7320.610687	7142.857143	3355.555556	5828.571429

```
In [20]: #Find the which is most demand occupation in male category
data.loc[data['Gender']=='Male']['Occupation'].value_counts(normalize=True) \
.plot(kind='bar',title="Visualize the Males Occupation",color=['#58FA58','#A9F5E1','#C9E8D8'])
plt.xlabel('Occupation')
plt.ylabel("Count of values")
plt.show()
```

### Visualize the Males Occupation




### Observations:

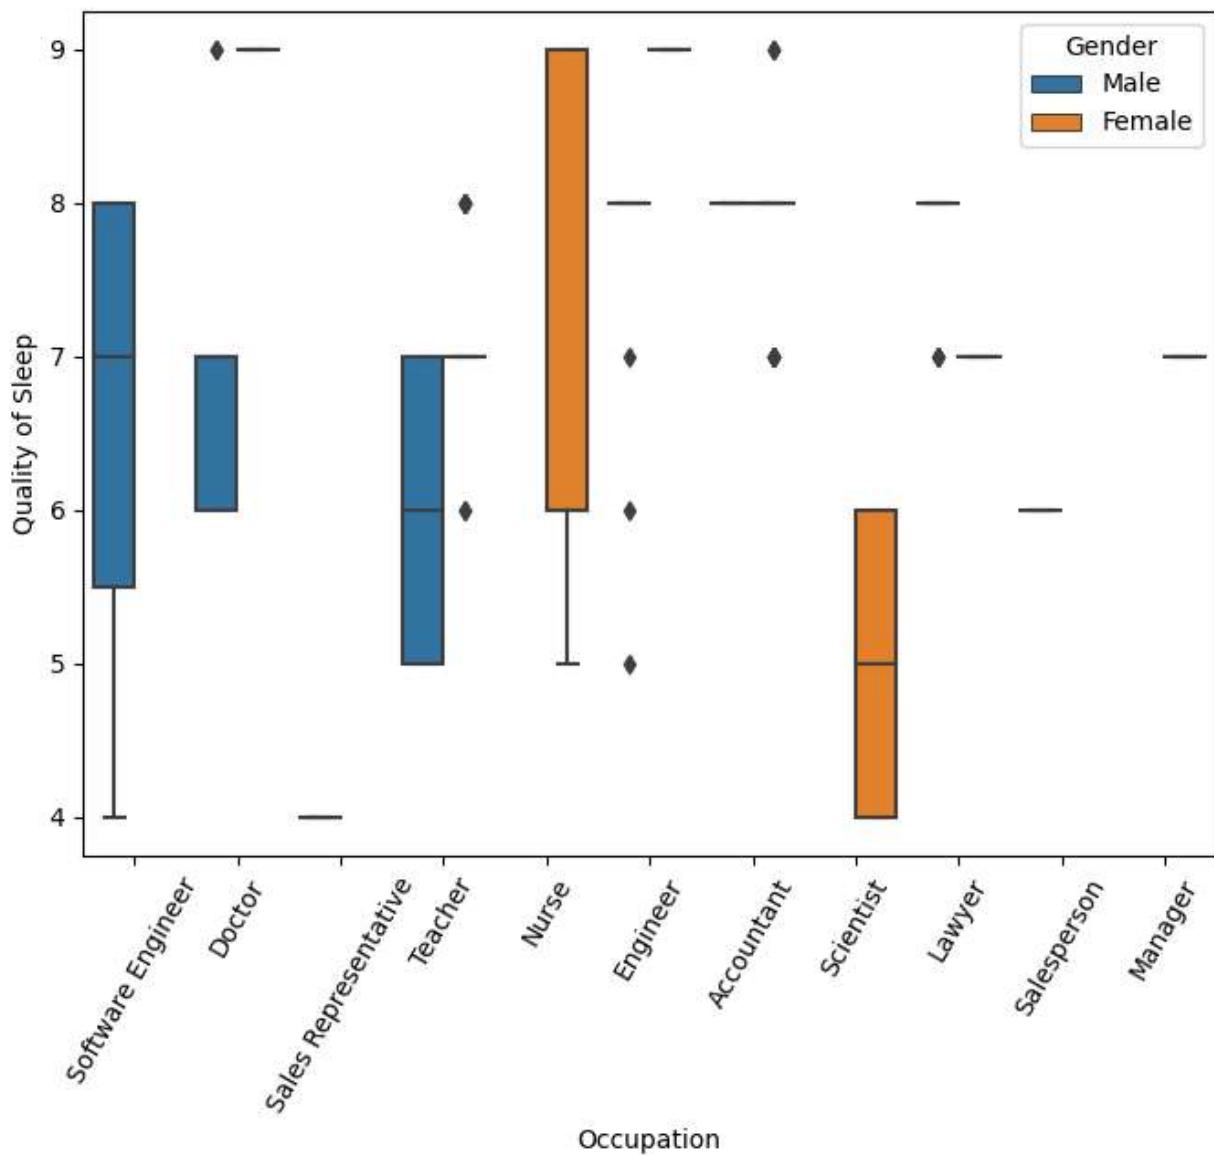
- From the above chart, we can observe an interesting pattern where overweight people tend to have higher stress levels, while normal-weight individuals have lower stress levels.
- Regarding the dominance of job roles in the male gender, it can be seen that the "Doctor" profession is the most dominating job, while the "Accountant" role is relatively less dominant.

```
In [21]: # Create a histogram for distributing the averages Daily steps
plt.figure(figsize=(8,6))
sns.histplot(data['Daily Steps'], bins='auto', color=['#2EEF7'], kde=True)
plt.title("Distribution of the Daily steps")
plt.show()
```

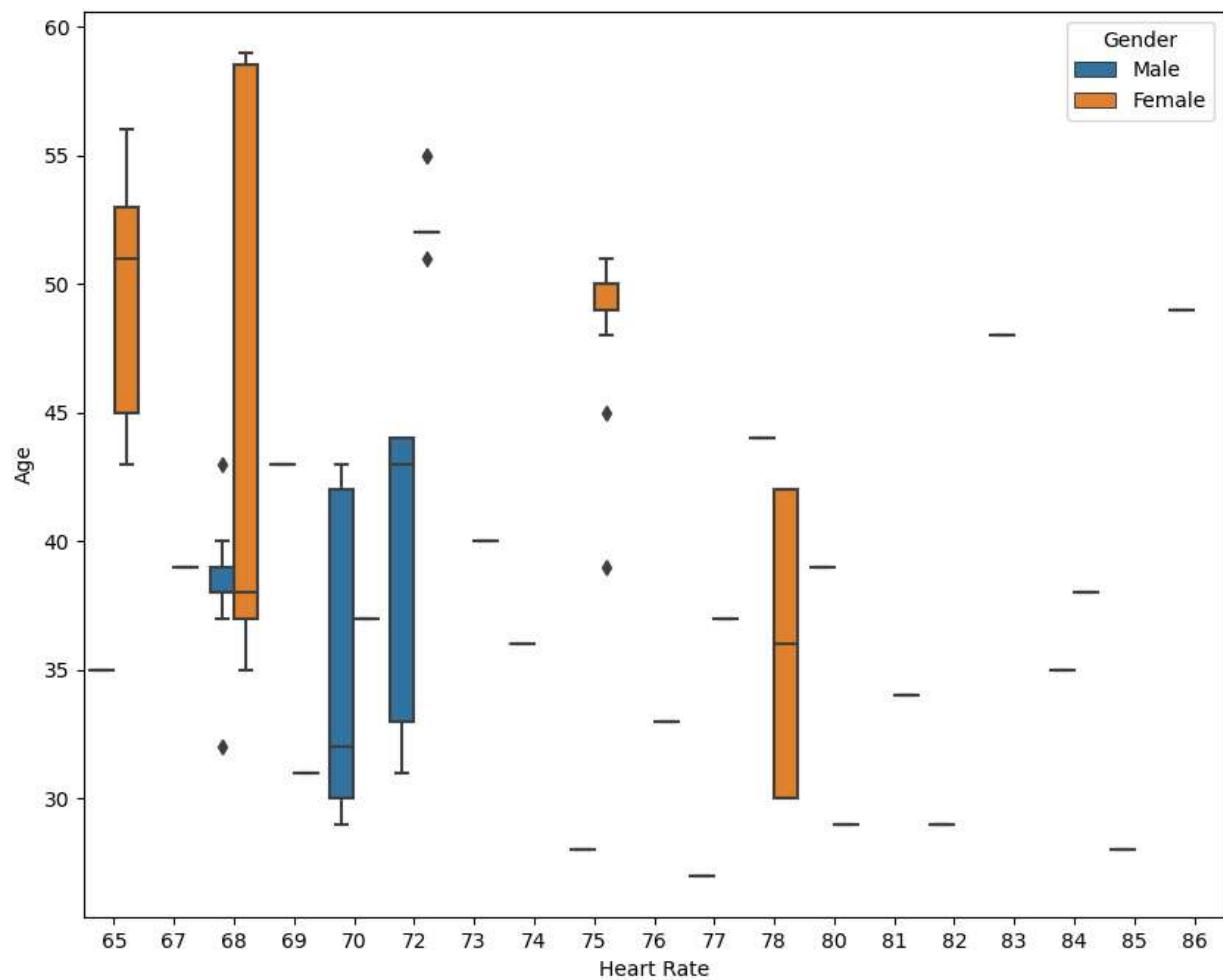
## Distribution of the Daily steps



```
In [22]: #Create a boxplot for occupation and quality of sleep on gender
plt.figure(figsize=(8,6))
sns.boxplot(data=data,x='Occupation',y='Quality of Sleep',hue='Gender',)
plt.xlabel("Occupation")
plt.ylabel("Quality of Sleep")
plt.xticks(rotation=60)
plt.show()
```

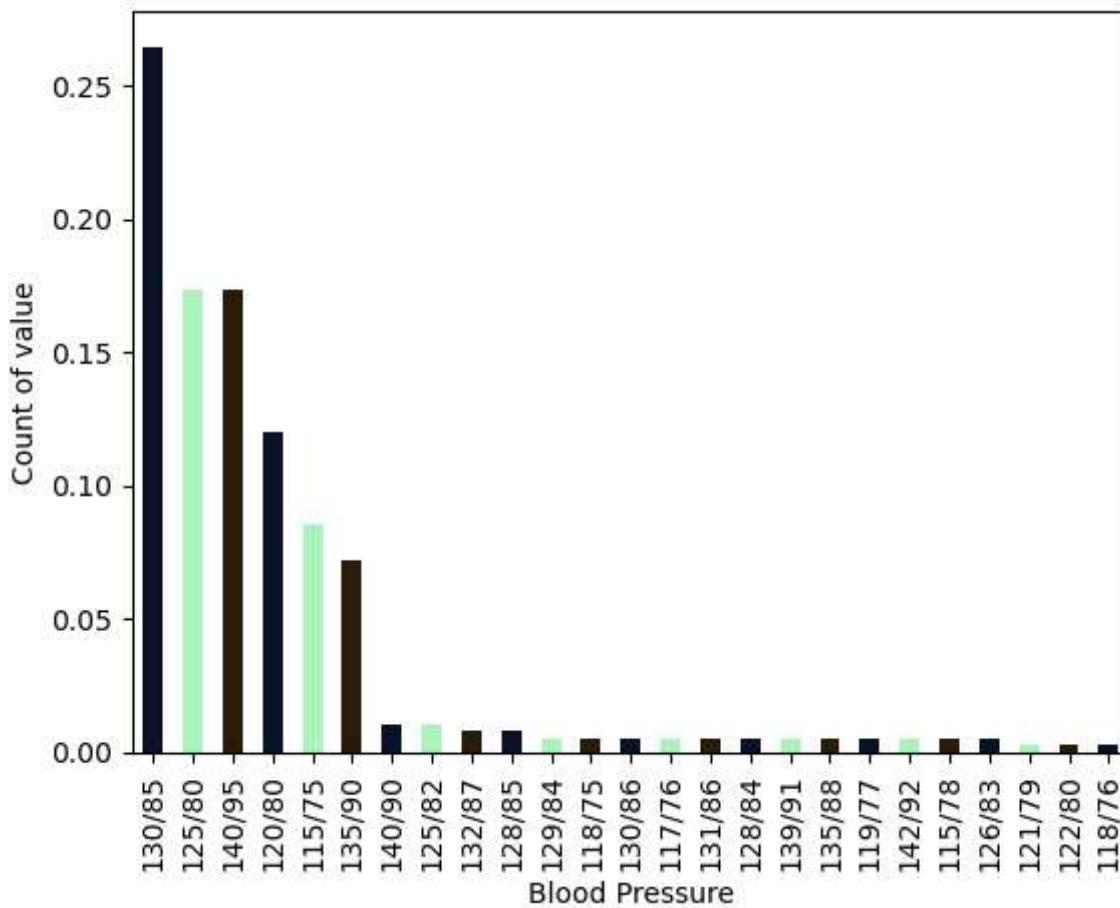


```
In [23]: #Visualize the boxplot with heart rate with age on gender
plt.figure(figsize=(10,8))
sns.boxplot(data=data,x='Heart Rate',y='Age',hue='Gender')
plt.xlabel("Heart Rate")
plt.ylabel("Age")
plt.show()
```



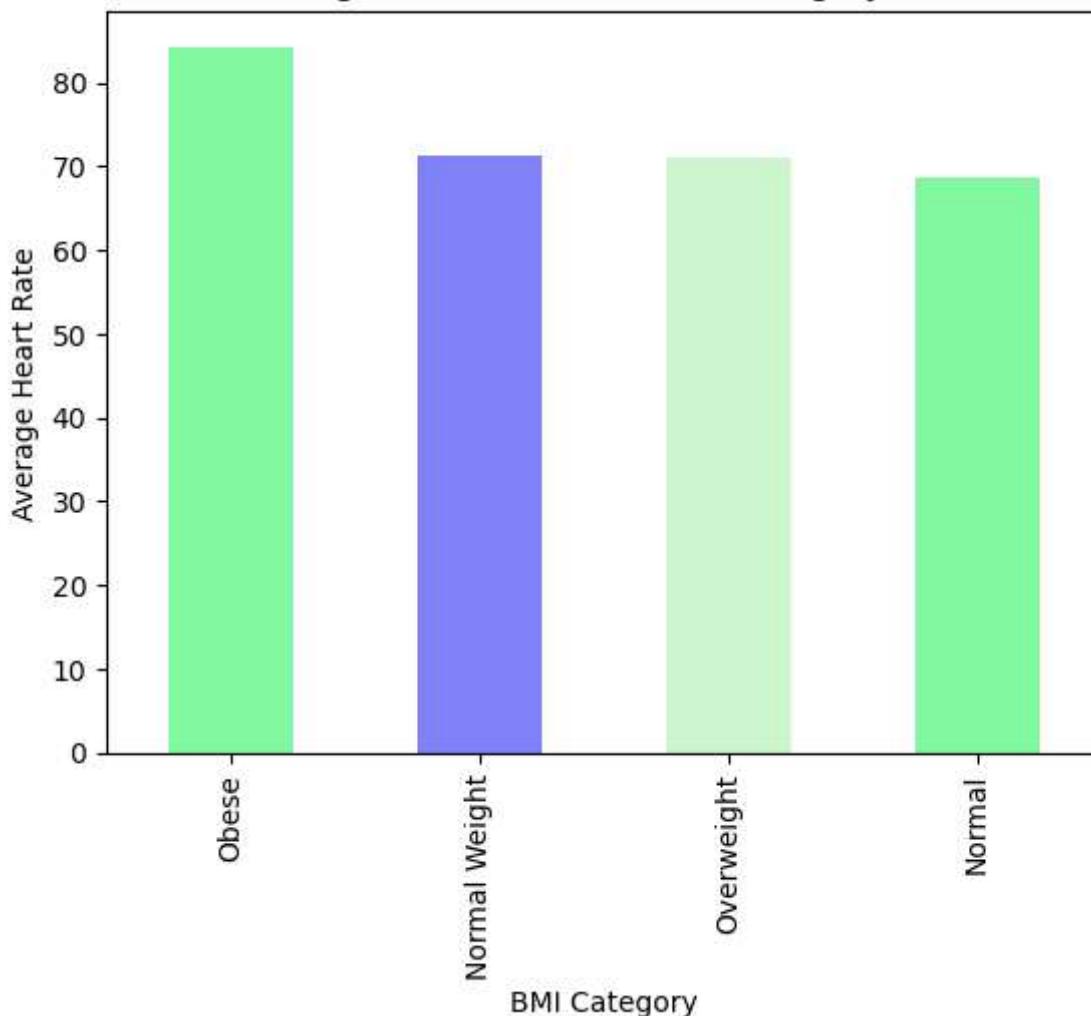
```
In [24]: #Create a bar chart with blood Pressure
data['Blood Pressure'].value_counts(normalize=True).sort_values(ascending=False) \
.plot(kind='bar',title="visualize the Blood Pressure",color=['#0A122A','#A9F5BC','#2A2A2A']) \
plt.xlabel("Blood Pressure") \
plt.ylabel("Count of value") \
plt.show()
```

## visualize the Blood Pressure



```
In [25]: #Find the Averate heart rate with Bmi category wise
data.groupby('BMI Category')['Heart Rate'].mean().sort_values(ascending=False) \
.plot(kind='bar',title='Average Heart Rate with BMI Category wise',color=['#81F79F','#
plt.xlabel("BMI Category")
plt.ylabel("Average Heart Rate")
plt.show()
```

### Average Heart Rate with BMI Category wise



```
In [26]: #Find the Average sleep duration with occupation with BMI Category on sleep Duration  
data.groupby(['Occupation','BMI Category'])['Sleep Duration'].mean().sort_values(ascending=True)  
.unstack()\n.style.background_gradient(cmap='gist_heat_r')
```


Out[26]:

BMI Category	Normal	Normal Weight	Obese	Overweight
Occupation				
Accountant	7.161538	7.520000	nan	6.566667
Doctor	6.887692	8.200000	7.700000	nan
Engineer	8.080357	7.450000	nan	6.966667
Lawyer	7.438095	6.900000	7.400000	7.100000
Manager	nan	nan	nan	6.900000
Nurse	nan	6.900000	nan	7.080303
Sales Representative	nan	nan	5.900000	nan
Salesperson	nan	nan	nan	6.403125
Scientist	nan	nan	nan	6.000000
Software Engineer	nan	7.500000	5.900000	6.100000
Teacher	7.166667	nan	6.300000	6.615152

In [27]:

```
#Create a countplot for occupation on the gender
plt.figure(figsize=(8,6))
sns.countplot(data=data,x='Occupation',hue='Gender')
plt.title("Count of the Occupation on Gender")
plt.xlabel("Occupation")
plt.ylabel("Count of values")
plt.show()
```

Count of the Occupation on Gender



## Machine Learning Modeling

```
In [28]: from sklearn.model_selection import train_test_split, GridSearchCV, ShuffleSplit
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.preprocessing import MinMaxScaler, LabelEncoder
from xgboost import XGBClassifier
from catboost import CatBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
```

```
In [29]: #Convert the Categorical Columns to into numerical values
for col in data.select_dtypes(include='object'):
    label_encoder=LabelEncoder()
    label_encoder.fit(data[col].unique())
    data[col]=label_encoder.transform(data[col])
```

```
In [30]: # Divided into the data into dependent and independent variables
X=data.drop(['Sleep Disorder'],axis=1)
y=data['Sleep Disorder']
#Scaling the values
scaler=MinMaxScaler()
scaler.fit(X)
X=scaler.fit_transform(X)
```

```
# Split the data into train and test
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.30,random_state=200)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(261, 11)
(113, 11)
(261,)
(113,)
```

```
In [31]: algo={

    'logistic_regression': {
        'model': LogisticRegression(),
        'params': {
        }
    },
    'SVM': {
        'model': SVC(),
        'params': {
        }
    },
    'decision_tree': {
        'model': DecisionTreeClassifier(),
        'params': {
            'criterion': ['gini', 'entropy', 'log_loss'],
            'splitter': ['best','random']
        }
    },
    'random_forest': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [1,5,10,20,50],
            'criterion': ['gini', 'entropy', 'log_loss']
        }
    },
    'ada_boost_classifier': {
        'model': AdaBoostClassifier(),
        'params': {
            'n_estimators': [1,5,10,20,50,100]
        }
    },
    'gradient_boosting_classifier': {
        'model': GradientBoostingClassifier(),
        'params': {
            'n_estimators': [1,5,10,20,50,100],
        }
    },
    'xgbclassifier': {
        'model': XGBClassifier(),
        'params': {
            'n_estimators': [1,5,10,20,50]
        }
    },
    'cat_boost_calssifier': {
        'model': CatBoostClassifier(),
        'params': {
            'iterations':[1]
        }
    }
}
```

```

        },
        'knn_classifier': {
            'model': KNeighborsClassifier(),
            'params': {
            }
        }
    }
}

```

In [32]:

```

score=[]
cv = ShuffleSplit(n_splits=5, test_size=0.2, random_state=10)
for algo_name,mp in algo.items():
    grid=GridSearchCV(mp['model'],mp['params'],cv=cv,return_train_score=False)
    grid.fit(X,y)
    score.append({
        'model':algo_name,
        'best_score':grid.best_score_,
        'best_params':grid.best_params_
    })

```

```

Learning rate set to 0.5
0:      learn: 0.6587861      total: 136ms      remaining: 0us
Learning rate set to 0.5
0:      learn: 0.6686615      total: 3.69ms      remaining: 0us
Learning rate set to 0.5
0:      learn: 0.6705082      total: 2.65ms      remaining: 0us
Learning rate set to 0.5
0:      learn: 0.6458583      total: 3.04ms      remaining: 0us
Learning rate set to 0.5
0:      learn: 0.6547457      total: 1.61ms      remaining: 0us
Learning rate set to 0.5
0:      learn: 0.6587434      total: 2.13ms      remaining: 0us

```

In [33]:

```
#Create a data frame for the
score=pd.DataFrame(score,columns=['model','best_score','best_params'])
score.head()
```

Out[33]:

	model	best_score	best_params
0	logistic_regression	0.877333	{}
1	SVM	0.882667	{}
2	decision_tree	0.874667	{'criterion': 'entropy', 'splitter': 'random'}
3	random_forest	0.896000	{'criterion': 'gini', 'n_estimators': 10}
4	ada_boost_classifier	0.858667	{'n_estimators': 5}

In [34]:

```
#Imbalanced the data set and create a function for the machine Learning model
from imblearn.combine import SMOTETomek
sm=SMOTETomek()
X_train_sm,y_train_sm=sm.fit_resample(X_train,y_train)
def model_buliding(model,X_train,X_test,y_train,y_test):
    model.fit(X_train_sm,y_train_sm)
    y_pred=model.predict(X_test)
    print(f'Accuracy score of the {model} is {accuracy_score(y_test,y_pred)*100:.2f}')
    print(classification_report(y_test,y_pred))
    confusion=confusion_matrix(y_test,y_pred)
    print(sns.heatmap(confusion,annot=True,cmap='gist_earth'))
```

In [35]: *#Applying the Catboostclassifier*

```
logistic=CatBoostClassifier(iterations=1)
model_buliding(logistic,X_train,X_test,y_train,y_test)
```

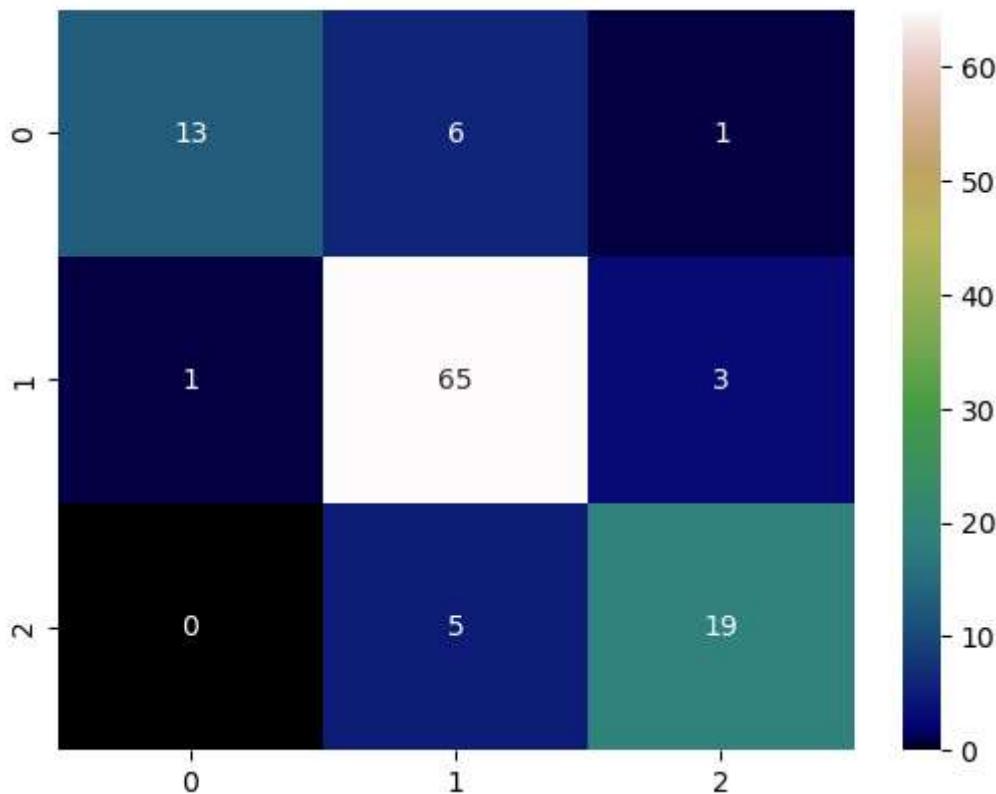
Learning rate set to 0.5

0: learn: 0.6650302 total: 2.71ms remaining: 0us

Accuracy score of the <catboost.core.CatBoostClassifier object at 0x000002569D3DF970> is 85.84

	precision	recall	f1-score	support
0	0.93	0.65	0.76	20
1	0.86	0.94	0.90	69
2	0.83	0.79	0.81	24
accuracy			0.86	113
macro avg	0.87	0.79	0.82	113
weighted avg	0.86	0.86	0.85	113

AxesSubplot(0.125,0.11;0.62x0.77)




In [36]: *#Applying the Catboostclassifier*

```
xgb=XGBClassifier()
model_buliding(xgb,X_train,X_test,y_train,y_test)
```

```
Accuracy score of the XGBClassifier(base_score=None, booster=None, callbacks=None,
colsample_bylevel=None, colsample_bynode=None,
colsample_bytree=None, early_stopping_rounds=None,
enable_categorical=False, eval_metric=None, feature_types=None,
gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
interaction_constraints=None, learning_rate=None, max_bin=None,
max_cat_threshold=None, max_cat_to_onehot=None,
max_delta_step=None, max_depth=None, max_leaves=None,
min_child_weight=None, missing=nan, monotone_constraints=None,
n_estimators=100, n_jobs=None, num_parallel_tree=None,
objective='multi:softprob', predictor=None, ...) is 92.04
```

	precision	recall	f1-score	support
0	0.89	0.85	0.87	20
1	0.97	0.94	0.96	69
2	0.81	0.92	0.86	24
accuracy			0.92	113
macro avg	0.89	0.90	0.90	113
weighted avg	0.92	0.92	0.92	113

AxesSubplot(0.125,0.11;0.62x0.77)




```
In [37]: #Applying the Catboostclassifier
knn=KNeighborsClassifier()
model_buliding(knn,X_train,X_test,y_train,y_test)
```

Accuracy score of the KNeighborsClassifier() is 88.50  
precision recall f1-score support

0	0.80	0.80	0.80	20
1	0.97	0.90	0.93	69
2	0.76	0.92	0.83	24

accuracy			0.88	113
macro avg	0.84	0.87	0.85	113
weighted avg	0.89	0.88	0.89	113

AxesSubplot(0.125,0.11;0.62x0.77)



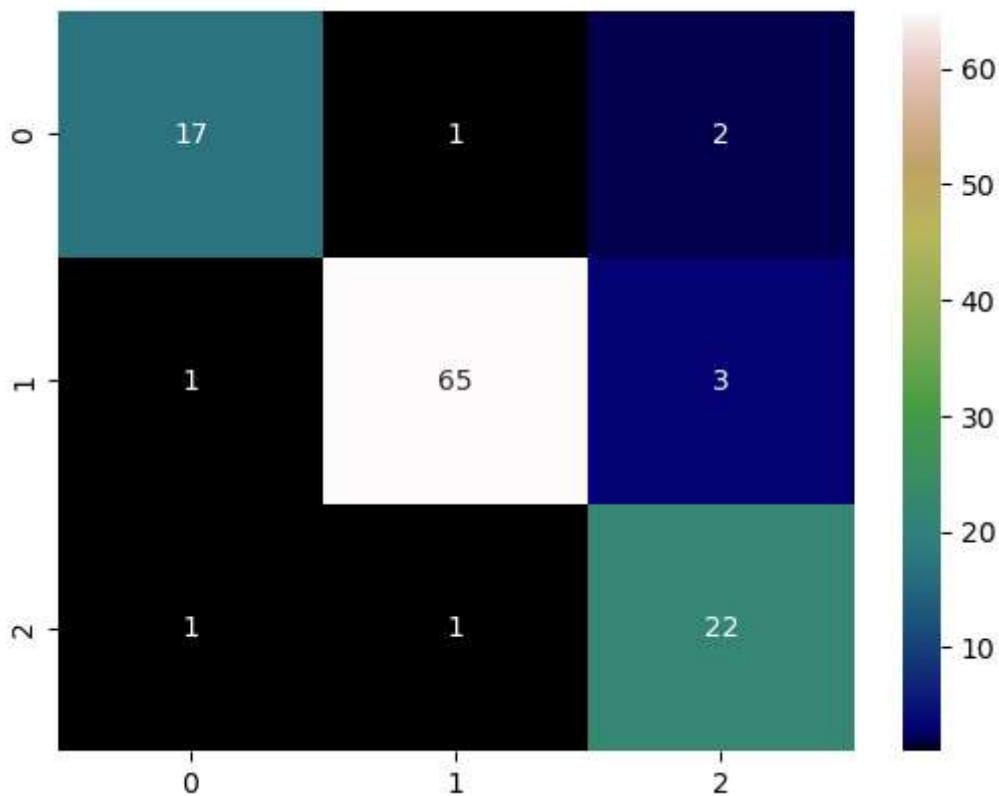
```
In [38]: #Applying the Catboostclassifier
grad=GradientBoostingClassifier()
model_buliding(grad,X_train,X_test,y_train,y_test)
```

Accuracy score of the GradientBoostingClassifier() is 92.04  
precision recall f1-score support

0	0.89	0.85	0.87	20
1	0.97	0.94	0.96	69
2	0.81	0.92	0.86	24

accuracy			0.92	113
macro avg	0.89	0.90	0.90	113
weighted avg	0.92	0.92	0.92	113

AxesSubplot(0.125,0.11;0.62x0.77)



## Conclusion:

This is another project in data science where data is obtained from Kaggle. The usual data preprocessing steps are performed, including data cleaning. Exploratory data analysis (EDA) techniques are applied to gain insights from the data, and questions are formulated based on the analysis. Visualizations such as bar charts, pie charts, and other types of charts are created to present the findings. Machine learning algorithms, specifically GradientBoostingClassifier and XGBClassifier, are applied to the data, and they achieve high accuracy scores.

In [ ]: